

Root Finding Methods

Chapter 7

Stats 102A: Introduction to Computational Statistics with R

UCLA



Acknowledgements: Miles Chen and Michael Tsiang

Introduction

Bisection Method

Fixed Point Iteration

Newton-Raphson Method

Secant Method

Introduction

Introduction

A **numerical method** is a technique or algorithm that approximates the solution to a mathematical problem. Numerical methods are typically iterative algorithms that are implemented using computers.

Why use approximate methods?

- ▶ The mathematical problem may not have an exact (or **analytical**) solution using algebra or calculus.
- ▶ Even if an exact solution exists in theory, it may be computationally **intractable**, i.e., it would take an unreasonably large amount of resources (usually time) to compute.

Numerical methods take advantage of a computer's ability to follow instructions quickly to compute a arbitrarily good approximations to solutions that might otherwise be impossible or infeasible to compute exactly.

Root Finding Methods

The first class of numerical methods we will discuss are **root finding** methods, which approximate the roots (or zeroes) of a function.

Suppose that $f : \mathbb{R} \rightarrow \mathbb{R}$ is a continuous function. A **root** of f is a solution to the equation

$$f(x) = 0.$$

Graphically, a root is where the graph $y = f(x)$ crosses the x -axis.

The solution to many mathematical and statistical problems can be expressed in terms of the root of a suitable function, so these methods have far reaching applications.

References and Resources

The main reference for these notes

- ▶ Jones, O., Maillardet, R., and Robinson, A.'s "Introduction to Scientific Programming and Simulation Using R":
<https://doi.org/10.1201/9781420068740>

Additional resources

- ▶ Jeffrey R. Chasnov's lecture notes on numerical methods:
<https://www.math.ust.hk/~machas/numerical-methods.pdf>
- ▶ Richard Anstee's notes on the Newton-Raphson method:
<https://www.math.ubc.ca/~anstee/math104/newtonmethod.pdf>

Bisection Method

Bisection Method

The **bisection method** is one of the simplest numerical methods for root-finding to implement, and it almost always works. The main disadvantage is that convergence is relatively slow.

The idea behind the bisection method is **root-bracketing**, which works by first finding an interval in which the root must lie, and then successively refining the bounding interval in such a way that the root is guaranteed to always lie inside the interval.

In the bisection method, the width of the bounding interval is successively halved.

Bisection Method

Suppose that f is a continuous function. Then f has a root in the interval (x_ℓ, x_r) if either:

- ▶ $f(x_\ell) < 0$ and $f(x_r) > 0$, or
- ▶ $f(x_\ell) > 0$ and $f(x_r) < 0$.

A convenient way to verify this condition is to check if

$$f(x_\ell)f(x_r) < 0.$$

If x_ℓ and x_r satisfy this condition, we say x_ℓ and x_r **bracket** the root.

The bisection method works by taking an interval (x_ℓ, x_r) that contains a root, then successively refining x_ℓ and x_r until $x_r - x_\ell \leq \varepsilon$, where $\varepsilon > 0$ is the pre-specified tolerance level.

Bisection Method

The **bisection** algorithm is described as follows:

Start with $x_\ell < x_r$ such that $f(x_\ell)f(x_r) < 0$.

1. If $|x_r - x_\ell| \leq \varepsilon$, then stop.
2. Compute $x_m = \frac{x_\ell + x_r}{2}$. If $f(x_m) = 0$, then stop.
3. If $f(x_\ell)f(x_m) < 0$, then assign $x_r = x_m$.
Otherwise, assign $x_\ell = x_m$.
4. Go back to step 1.

Notice that, at every iteration, the interval (x_ℓ, x_r) always contains a root. Assuming we start with $f(x_\ell)f(x_r) < 0$, the algorithm is guaranteed to converge, with the approximation error reducing by a constant factor $1/2$ at each iteration. If we stop when $x_r - x_\ell \leq \varepsilon$, then we know that both x_ℓ and x_r are within ε distance of a root.

Question: When will the bisection method fail to find a root?

Example of Bisection Method

An implementation of bisection method for $f(x) = \log(x) - e^{-x}$:

```
f <- function(x){  
  log(x) - exp(-x)  
}
```

```
x_l <- 1  
x_r <- 2  
tol <- 1e-8  
f_l <- f(x_l)  
f_r <- f(x_r)  
f_l * f_r < 0
```

```
## [1] TRUE
```

Example of Bisection Method

```
its <- 0
while(x_r - x_l > tol){
  x_m <- (x_l + x_r) / 2
  f_m <- f(x_m)
  if(identical(all.equal(f_m,0),TRUE)){break}
  if(f_l * f_m < 0){
    x_r <- x_m
  } else{
    x_l <- x_m
  }
  its <- its + 1
}
x_m
```

```
## [1] 1.3098
```

```
its
```

```
## [1] 24
```

Fixed Point Iteration

Fixed Points

Let $g : \mathbb{R} \rightarrow \mathbb{R}$ be a continuous function. A **fixed point** of g is a solution to the equation

$$g(x) = x.$$

Graphically, a fixed point is where the graph $y = g(x)$ crosses the line $y = x$.

An interesting/useful property of fixed points is that they remain fixed after iteratively applying the function an arbitrary number of times. That is, if x is a fixed point of g , then

$$x = g(x) = g(g(x)) = g(g(g(x))) = \cdots.$$

Fixed Points for Root Finding

The problem of finding fixed points can be framed as a root-finding problem. Define the function $f(x)$ by the equation

$$f(x) = c[g(x) - x],$$

for some constant $c \neq 0$. Then $f(x) = 0$ if and only if $g(x) = x$.

Conversely, the problem of finding solutions to $f(x) = 0$ is equivalent to finding the fixed points of the function

$$g(x) = c \cdot f(x) + x.$$

Note: There are often many ways to write $f(x) = 0$ as $g(x) = x$, and, in practice, some are better than others.

Fixed Point Iteration

The **fixed point iteration** algorithm is described as follows:

- ▶ Start from an initial value x_0 .
- ▶ Compute $g(x_0)$ as the next value x_1 .
- ▶ Repeat:

$$x_{n+1} = g(x_n)$$

To avoid iterating forever, we stop when $|x_n - x_{n-1}| \leq \varepsilon$, for some user-specified **tolerance** level $\varepsilon > 0$.

Example of Fixed Point Iteration

Suppose we want to find the solution to the equation

$$f(x) = \log(x) - e^{-x} = 0.$$

We first want to express this as a fixed point problem.

One way to rewrite $f(x)$ is to solve for x :

$$\begin{aligned}\log(x) - e^{-x} &= 0 \\ \log(x) &= e^{-x} \\ x &= e^{e^{-x}}.\end{aligned}$$

Let $g(x) = e^{e^{-x}}$. Then the fixed point of $g(x)$ corresponds to the root of $f(x)$.

Example of Fixed Point Iteration

An implementation of fixed point iteration for $g(x) = e^{e^{-x}}$:

```
g <- function(x){  
  exp(exp(-x))  
}
```

```
x_old <- 10  
x_new <- g(x_old)  
its <- 0  
tol <- 1e-8  
while(abs(x_new - x_old) > tol){  
  x_old <- x_new  
  x_new <- g(x_old)  
  its <- its + 1  
}
```

Note: For algorithms that might not always converge (or are possibly slow to converge), it is good practice to have an option for stopping the algorithm after a specified maximum number of iterations. This prevents your code from running forever.

Example of Fixed Point Iteration

Our results:

```
x_new
```

```
## [1] 1.3098
```

```
log(x_new) - exp(-x_new)
```

```
## [1] 9.711432e-10
```

```
its
```

```
## [1] 19
```

Fixed Point Convergence

Suppose that the fixed point iteration method converges, i.e., $x_{n+1} = g(x_n)$ and $x_n \rightarrow a$, for some a .

Given g is continuous, then

$$a = \lim_{n \rightarrow \infty} x_{n+1} = \lim_{n \rightarrow \infty} g(x_n) = g\left(\lim_{n \rightarrow \infty} x_n\right) = g(a),$$

so a is a fixed point of g .

That is, *if* the fixed point iteration method converges, then it converges to a fixed point.

But does fixed point iteration always converge?

Fixed Point Convergence

The convergence of fixed point iteration relies on the existence of an attractive fixed point. If the fixed point is attractive, and the initial value is relatively “close” to the fixed point, then the fixed point iteration method is guaranteed to converge.

Warning: Not all functions have fixed points, and not all fixed points are attractive. Thus, in general, the fixed point iteration method is not guaranteed to converge.

A sufficient condition depends on the value of the *derivative* $g'(x)$ at the fixed point: If a is a fixed point of g and $|g'(a)| < 1$, then a is attractive.¹

If $|g'(a)| > 1$, then the fixed point is called **repelling**, and the fixed point iteration method will diverge.

¹The derivative condition $|g'(a)| < 1$ is often the usual definition for an attractive fixed point, and the convergence of fixed point iteration is a consequence. For more information:

<http://math.mit.edu/classes/18.01/F2011/lecture3.pdf>

Newton-Raphson Method

Newton-Raphson Method

Fixed point iteration is considered relatively slow, as the approximation error is usually divided by a constant factor at each iteration.

A well-known method for much faster convergence that is still widely used today is the **Newton-Raphson** method.

The main idea behind Newton-Raphson is to use the property that the tangent line at a point is the best linear approximation to the function near that point.

Linear Approximation

Let a be a root of f . Suppose that f is differentiable with continuous derivative f' .

Let x_0 denote the initial guess at a . Consider the Taylor series expansion of $f(x)$ around $x = x_0$, given by

$$f(x) = f(x_0) + f'(x_0)(x - x_0) + \frac{f''(x_0)}{2}(x - x_0)^2 + \cdots .$$

By dropping higher order terms, the **linear approximation** to $f(x)$ close to x_0 is the equation of the tangent line

$$f(x) \approx f(x_0) + f'(x_0)(x - x_0).$$

Linear Approximation

Since the tangent line is a good approximation to f near x_0 , the next guess x_1 is the point where the tangent line crosses the x -axis.

That is, we choose x_1 such that

$$0 = f(x_0) + f'(x_0)(x_1 - x_0)$$

$$0 = f(x_0) + f'(x_0)x_1 - f'(x_0)x_0$$

$$f'(x_0)x_1 = f'(x_0)x_0 - f(x_0)$$

$$x_1 = x_0 - \frac{f(x_0)}{f'(x_0)}.$$

Newton-Raphson Method

Repeating the same procedure, the next guess x_2 is then

$$x_2 = x_1 - \frac{f(x_1)}{f'(x_1)},$$

and, in general,

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}.$$

Newton-Raphson Method

The **Newton-Raphson** algorithm is described as follows:

- ▶ Start from an initial value x_0 .
- ▶ Repeat:

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

Stop when $|f(x_n)| \leq \varepsilon$, where $\varepsilon > 0$ is the pre-specified tolerance level.

Example of Newton-Raphson Method

Suppose we want to use the Newton-Raphson method to find the root of

$$f(x) = \log(x) - e^{-x}.$$

Taking the derivative with respect to x , we have

$$f'(x) = \frac{1}{x} + e^{-x}.$$

Example of Newton-Raphson Method

An implementation of Newton-Raphson for $f(x) = \log(x) - e^{-x}$:

```
f <- function(x){  
  log(x) - exp(-x)  
}  
f_prime <- function(x){  
  (1/x) + exp(-x)  
}
```

```
x <- 2  
its <- 0  
tol <- 1e-8  
while(abs(f(x)) > tol){  
  x <- x - f(x) / f_prime(x)  
  its <- its + 1  
}
```

Example of Newton-Raphson Method

```
x
```

```
## [1] 1.3098
```

```
log(x) - exp(-x)
```

```
## [1] -3.494008e-09
```

```
its
```

```
## [1] 4
```

Newton-Raphson Convergence

Suppose that the Newton-Raphson method converges, i.e., $x_n \rightarrow a$, for some a . Given f is differentiable and f' is continuous, then

$$\begin{aligned} a &= \lim_{n \rightarrow \infty} x_{n+1} = \lim_{n \rightarrow \infty} \left[x_n - \frac{f(x_n)}{f'(x_n)} \right] \\ &= \lim_{n \rightarrow \infty} x_n - \frac{\lim_{n \rightarrow \infty} f(x_n)}{\lim_{n \rightarrow \infty} f'(x_n)} \\ &= a - \frac{f(a)}{f'(a)}. \end{aligned}$$

Thus, assuming $f'(a) \neq \pm\infty$, then $f(a) = 0$ (i.e., a is a root of f).

That is, *if* the Newton-Raphson method converges, and the derivative at a is finite, then it converges to a root.

But does Newton-Raphson always converge?

Newton-Raphson Convergence

It can be shown that if f is “well behaved” at a (i.e., $f'(a) \neq 0$ and f'' is finite and continuous at a), and you start with an initial value x_0 that is “close enough” to a , then x_n will converge quickly.

In fact, in most cases, the Newton-Raphson method converges **quadratically**, meaning the number of correct digits doubles per iteration.

However, when the conditions for quadratic convergence are not satisfied, then the Newton-Raphson method is not guaranteed to converge as quickly, and it may not even converge at all.

Side Note: A proof and more exact conditions for quadratic convergence can be found here:

https://en.wikipedia.org/wiki/Newton%27s_method#Analysis

Newton-Raphson Convergence

The Newton-Raphson method can fail for several reasons.

- ▶ Convergence can depend on the initial guess. If the starting value is too far from the root, the method may either not converge or converge to a different root.
- ▶ If $f'(x_n) = 0$ for any n , then the tangent line is horizontal and will not cross the x -axis. There can be issues even if $f'(x_n)$ is close to 0.
- ▶ If guesses oscillate back and forth in a cycle, then the Newton-Raphson method will not converge.
- ▶ The second derivative $f''(x)$ is infinite or undefined at the root.
- ▶ In some non-ideal scenarios (e.g., when $f'(a) = 0$), the Newton-Raphson method may still converge, but convergence may be much slower than quadratic.

Animated examples: <https://yihui.name/animation/example/newton-method/>

Side Note: A brief summary of the pros and cons (mostly the cons) of Newton-Raphson method is shown here:

http://www.cas.mcmaster.ca/~cs4te3/notes/newtons_method.pdf

Secant Method

Secant Method

Because of its fast convergence, the Newton-Raphson method is often the preferred method for root-finding when applicable.

One immediate downside to Newton-Raphson, however, is that it requires the analytic computation of the derivative of $f(x)$.

If the derivative is difficult to compute or does not exist, then we need a different method. An often used substitute method is the **secant method**.

Like the Newton-Raphson method, the secant method is based on a linear approximation to the function of interest.

Linear Approximation with the Secant Line

Let a be a root of f . We start with *two* initial guesses, x_0 and x_1 , for the value of a .

Consider the straight line which connects the points $(x_0, f(x_0))$ and $(x_1, f(x_1))$. This straight line is called a **secant line** to the curve $y = f(x)$. In a neighborhood around x_0 and x_1 , the secant line is a linear approximation to $f(x)$.

The slope of the secant line is

$$m = \frac{f(x_1) - f(x_0)}{x_1 - x_0},$$

so the equation of the secant line, in point-slope form, is

$$y - f(x_1) = m(x - x_1) = \frac{f(x_1) - f(x_0)}{x_1 - x_0}(x - x_1),$$

or, in slope-intercept form,

$$y = \frac{f(x_1) - f(x_0)}{x_1 - x_0}(x - x_1) + f(x_1).$$

Linear Approximation with the Secant Line

Since the secant line is a good approximation to f near x_1 , the next guess x_2 is then the point where the secant line crosses the x -axis.

That is, we choose x_2 such that

$$\begin{aligned}0 &= \frac{f(x_1) - f(x_0)}{x_1 - x_0}(x_2 - x_1) + f(x_1) \\ \frac{f(x_1) - f(x_0)}{x_1 - x_0}x_2 &= \frac{f(x_1) - f(x_0)}{x_1 - x_0}x_1 - f(x_1) \\ x_2 &= x_1 - f(x_1)\frac{x_1 - x_0}{f(x_1) - f(x_0)}.\end{aligned}$$

Secant Method

Repeating the same procedure, the next guess x_3 is then

$$x_3 = x_2 - f(x_2) \frac{x_2 - x_1}{f(x_2) - f(x_1)},$$

and, in general,

$$x_{n+1} = x_n - f(x_n) \frac{x_n - x_{n-1}}{f(x_n) - f(x_{n-1})}.$$

Note: Notice that if x_n and x_{n-1} are close to each other, then

$$f'(x_n) \approx \frac{f(x_n) - f(x_{n-1})}{x_n - x_{n-1}},$$

which is why the secant method is often seen as an approximation to the Newton-Raphson method, where we approximate the derivative by a finite difference approximation.

Secant Method

The **secant method** algorithm is described as follows:

- ▶ Start from initial values x_0 and x_1 .
- ▶ Repeat:

$$x_{n+1} = x_n - f(x_n) \frac{x_n - x_{n-1}}{f(x_n) - f(x_{n-1})}$$

Stop when $|f(x_n)| \leq \varepsilon$, where $\varepsilon > 0$ is the pre-specified tolerance level.

Secant Method Convergence

The convergence properties of the secant method are similar to those of the Newton-Raphson method.

If f is well behaved at a , and you start with initial values x_0 and x_1 sufficiently close to a , then x_n will converge to a quickly, though not quite as fast as the Newton-Raphson method.

Just as for the Newton-Raphson method, we cannot guarantee convergence for the secant method.

The main trade-offs in using the secant method over Newton-Raphson: We no longer need to know f' , but we lose some speed and have to provide two initial points, x_0 and x_1 .

Best of Both Worlds

The most popular current root-finding methods use root-bracketing to get close to a root, then switch over to the Newton-Raphson or secant method when it seems safe to do so.

This strategy combines the safety of bisection with the speed of the secant method.

Order of Convergence

Let a be the root and x_n be the n th approximation to the root. Define the **error** to be

$$\varepsilon_n = a - x_n.$$

A root finding method has an **order of convergence** p if, for large n , we have the approximate relationship

$$|\varepsilon_{n+1}| = k|\varepsilon_n|^p,$$

for some positive constant k . Larger values of p correspond to faster convergence.

Side Note: The derivations for the orders of convergence discussed here are in Jeffrey R. Chasnov's lecture notes on numerical methods:
<https://www.math.ust.hk/~machas/numerical-methods.pdf>

Order of Convergence

The order of convergence of the bisection method is 1, since the error is reduced by approximately a factor of 2 with each iteration. That is,

$$|\varepsilon_{n+1}| = \frac{1}{2}|\varepsilon_n|.$$

Order of convergence 1 is also called **linear convergence**.

It can be shown that, under certain conditions, the Newton-Raphson method has order of convergence 2, i.e.,

$$|\varepsilon_{n+1}| = k|\varepsilon_n|^2,$$

which is also called **quadratic convergence**.

The secant method can be shown to have an order of convergence equal to

$$p = \frac{1 + \sqrt{5}}{2} \approx 1.618034,$$

which is the famous **golden ratio** (sometimes denoted by Φ)!